

assignment01-regression

March 27, 2017

1 CSAL4243: Introduction to Machine Learning

Muhammad Mudassir Khan (mudassir.khan@ucp.edu.pk)

2 Assignment 1: Linear Regression

In this assignment you are going to learn how Linear Regression works by using the code for linear regression and gradient descent we have been looking at in the class. You are also going to use linear regression from [scikit-learn](#) library for machine learning. You are going to learn how to download data from [kaggle](#) (a website for datasets and machine learning) and upload submissions to kaggle competitions. And you will be able to compete with the world.

2.0.1 Overview

- Pseudocode
- Tasks
 - Load and analyze data
- Task 1: Effect of Learning Rate α
 - Load X and y
 - Linear Regression with Gradient Descent code
 - Run Gradient Descent on training data
 - Plot trained line on data
- Task 2: Predict test data output and submit it to Kaggle
 - Upload .csv file to Kaggle.com
- Task 3: Use scikit-learn for Linear Regression
- Task 4: Multivariate Linear Regression
- Resources
- Credits

3 Pseudocode

3.1 Linear Regression with Gradient Descent

- Load training data into X_{train} and y_{train}
- [Optionally] normalize features X_{train} using $x^i = \frac{x^i - \mu^i}{\rho^i}$ where μ^i is mean and ρ^i is standard deviation of feature i
- Initialize hyperparameters
 - iterations
 - learning rate α
- Initialize θ_s
- At each iteration
 - Compute cost using $J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h(x^i) - y^i)^2$ where $h(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n$
 - Update θ_s using
 - [Optionally] Break if cost $J(\theta)$ does not change.

4 Tasks

1. Effect of Learning Rate α
2. Predict test data output and submit it to Kaggle
3. Use scikit-learn for Linear Regression
4. Multivariate Linear Regression

4.1 Load and analyze data

```
In [4]: %matplotlib inline
import pandas as pd
import numpy as np
import seaborn as sns
from sklearn import linear_model
import matplotlib.pyplot as plt
import matplotlib as mpl

# read house_train.csv data in pandas dataframe df_train using pandas read_
df_train = pd.read_csv('datasets/house_price/train.csv', encoding='utf-8')

In [5]: # check data by printing first few rows
df_train.head()
```

```
Out[5]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	\
0	1	60	RL	65.0	8450	Pave	NaN	Reg	
1	2	20	RL	80.0	9600	Pave	NaN	Reg	
2	3	60	RL	68.0	11250	Pave	NaN	IR1	
3	4	70	RL	60.0	9550	Pave	NaN	IR1	
4	5	60	RL	84.0	14260	Pave	NaN	IR1	

	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal
0	Lvl	AllPub	...	0	NaN	NaN	NaN	
1	Lvl	AllPub	...	0	NaN	NaN	NaN	
2	Lvl	AllPub	...	0	NaN	NaN	NaN	
3	Lvl	AllPub	...	0	NaN	NaN	NaN	
4	Lvl	AllPub	...	0	NaN	NaN	NaN	


	MoSold	YrSold	SaleType	SaleCondition	SalePrice
0	2	2008	WD	Normal	208500
1	5	2007	WD	Normal	181500
2	9	2008	WD	Normal	223500
3	2	2006	WD	Abnorml	140000
4	12	2008	WD	Normal	250000

[5 rows x 81 columns]


```
In [6]: # check columns in dataset
df_train.columns
```

```
Out[6]: Index(['Id', 'MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street',
               'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig',
               'LandSlope', 'Neighborhood', 'Condition1', 'Condition2', 'BldgType',
               'HouseStyle', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAd',
               'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType',
               'MasVnrArea', 'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual',
               'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1',
               'BsmtFinType2', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating',
               'HeatingQC', 'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF',
               'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBa',
               'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual',
               'TotRmsAbvGrd', 'Functional', 'Fireplaces', 'FireplaceQu', 'GarageTy',
               'GarageYrBlt', 'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQu',
               'GarageCond', 'PavedDrive', 'WoodDeckSF', 'OpenPorchSF',
               'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'PoolQC',
               'Fence', 'MiscFeature', 'MiscVal', 'MoSold', 'YrSold', 'SaleType',
               'SaleCondition', 'SalePrice'],
              dtype='object')
```

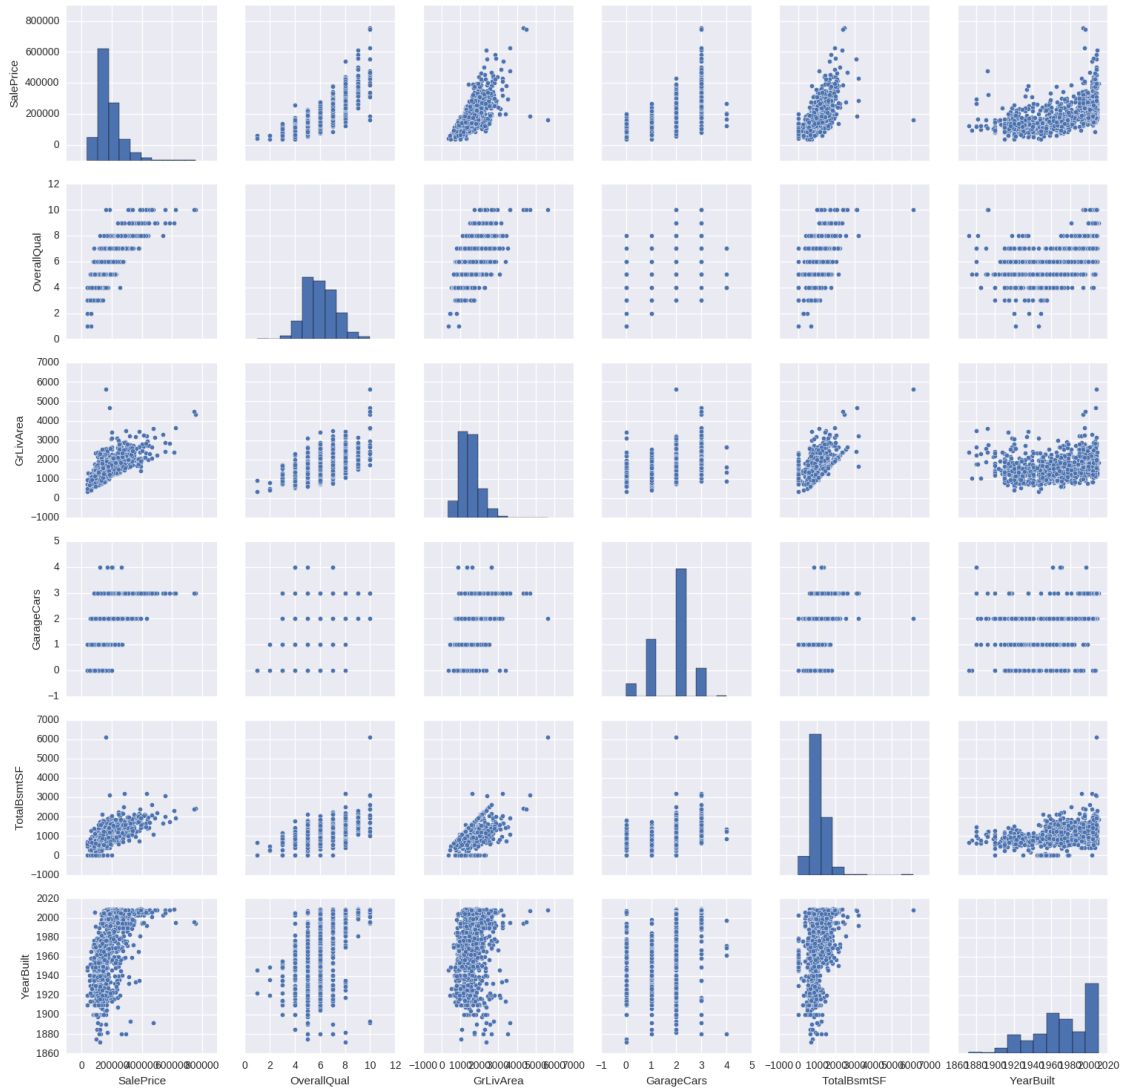
```
In [7]: # check correlation matrix, darker means more correlation
corrmat = df_train.corr()
f, ax_train= plt.subplots(figsize=(12, 9))
sns.heatmap(corrmat, vmax=.8, square=True);
```



```
In [8]: # SalePrice correlation matrix with top k variables
k = 10 #number of variables for heatmap
cols = corrmat.nlargest(k, 'SalePrice')['SalePrice'].index
cm = np.corrcoef(df_train[cols].values.T)
sns.set(font_scale=1.25)
hm = sns.heatmap(cm, cbar=True, annot=True, square=True, fmt='.2f', annot_k=
plt.show())
```



```
In [9]: #scatterplot with some important variables
cols = ['SalePrice', 'OverallQual', 'GrLivArea', 'GarageCars', 'TotalBsmtSF', '1stFlrSF', 'FullBath', 'TotRmsAbvGrd', 'YearBuilt']
sns.set()
sns.pairplot(df_train[cols], size = 2.5)
plt.show();
```



Task 1: Effect of Learning Rate α Use Linear Regression code below using X="GrLivArea" as input variable and y="SalePrice" as target variable. Use different values of α given in table below and comment on why they are useful or not and which one is a good choice.

- $\alpha = 0.000001$:
- $\alpha = 0.00000001$:
- $\alpha = 0.000000001$:

Load X and y

```
In [10]: # Load X and y variables from pandas dataframe df_train
cols = ['GrLivArea']
X_train = np.array(df_train[cols])
y_train = np.array(df_train[["SalePrice"]])
```

```

# Get m = number of samples and n = number of features
m = X_train.shape[0]
n = X_train.shape[1]

# append a column of 1's to X for theta_0
X_train = np.insert(X_train,0,1,axis=1)

```

4.2 Linear Regression with Gradient Descent code

```

In [11]: iterations = 1500
alpha = 0.000000001 # change it and find what happens

def h(X, theta): #Linear hypothesis function
    hx = np.dot(X,theta)
    return hx

def computeCost(theta,X,y): #Cost function
    """
    theta is an n- dimensional vector, X is matrix with n- columns and m-
    y is a matrix with m- rows and 1 column
    """
    #note to self: *.shape is (rows, columns)
    return float((1./(2*m)) * np.dot((h(X,theta)-y).T, (h(X,theta)-y)))


#Actual gradient descent minimizing routine
def gradientDescent(X,y, theta_start = np.zeros((n+1,1))):
    """
    theta_start is an n- dimensional vector of initial theta guess
    X is input variable matrix with n- columns and m- rows. y is a matrix
    """
    theta = theta_start
    j_history = [] #Used to plot cost as function of iteration
    theta_history = [] #Used to visualize the minimization path later on
    for meaninglessvariable in range(iterations):
        tmptheta = theta
        # append for plotting
        j_history.append(computeCost(theta,X,y))
        theta_history.append(list(theta[:,0]))
        #Simultaneously updating theta values
        for j in range(len(tmptheta)):
            tmptheta[j] = theta[j] - (alpha/m)*np.sum((h(X,theta) - y)*np.
        theta = tmptheta
    return theta, theta_history, j_history

```

4.3 Run Gradient Descent on training data

```
In [12]: #Actually run gradient descent to get the best-fit theta values
initial_theta = np.zeros((n+1,1));
theta, theta_history, j_history = gradientDescent(X_train,y_train,initial_theta)

plt.plot(j_history)
plt.title("Convergence of Cost Function")
plt.xlabel("Iteration number")
plt.ylabel("Cost function")
plt.show()
```



4.4 Plot trained line on data

```
In [13]: # predict output for training data
hx_train= h(X_train, theta)

# plot it
plt.scatter(X_train[:,1],y_train)
plt.plot(X_train[:,1],hx_train[:,0], color='red')
plt.show()
```




Task 2: Predict test data output and submit it to Kaggle In this task we will use the model trained above to predict “SalePrice” on test data. Test data has all the input variables/features but no target variable. Our aim is to use the trained model to predict the target variable for test data. This is called generalization i.e. how good your model works on unseen data. The output in the form “Id”, “SalePrice” in a .csv file should be submitted to kaggle. Please provide your score on kaggle after this step as an image. It will be compared to the 5 feature Linear Regression later.

```
In [14]: # read data in pandas frame df_test and check first few rows
         # write code here
```

```
df_test.head()
```

```
Out[14]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape
0	1461	20	RH	80.0	11622	Pave	NaN	Reg
1	1462	20	RL	81.0	14267	Pave	NaN	IR1
2	1463	60	RL	74.0	13830	Pave	NaN	IR1
3	1464	60	RL	78.0	9978	Pave	NaN	IR1
4	1465	120	RL	43.0	5005	Pave	NaN	IR1

	LandContour	Utilities	...	ScreenPorch	PoolArea	PoolQC	Fence
0	Lvl	AllPub	...	120	0	NaN	MnPrv
1	Lvl	AllPub	...	0	0	NaN	NaN
2	Lvl	AllPub	...	0	0	NaN	MnPrv
3	Lvl	AllPub	...	0	0	NaN	NaN
4	HLS	AllPub	...	144	0	NaN	NaN

	MiscFeature	MiscVal	MoSold	YrSold	SaleType	SaleCondition
0	NaN	0	6	2010	WD	Normal
1	Gar2	12500	6	2010	WD	Normal
2	NaN	0	3	2010	WD	Normal
3	NaN	0	6	2010	WD	Normal
4	NaN	0	1	2010	WD	Normal

[5 rows x 80 columns]

```
In [15]: # check statistics of test data, make sure no data is missing.
print(df_test.shape)
df_test[cols].describe()
```

(1459, 80)

```
Out[15]:
```

	GrLivArea
count	1459.000000
mean	1486.045922
std	485.566099
min	407.000000
25%	1117.500000
50%	1432.000000
75%	1721.000000
max	5095.000000

```
In [16]: # Get X_test, no target variable (SalePrice) provided in test data. It is
X_test = np.array(df_test[cols])
```

```
#Insert the usual column of 1's into the "X" matrix
X_test = np.insert(X_test,0,1,axis=1)
```

```
In [17]: # predict test data labels i.e. y_test
predict = h(X_test, theta)
```

```
In [18]: # save prediction as .csv file
pd.DataFrame({'Id': df_test.Id, 'SalePrice': predict[:,0]}).to_csv("predict1.csv")
```

4.5 Upload .csv file to Kaggle.com

- Create an account at <https://www.kaggle.com>
- Go to <https://www.kaggle.com/c/house-prices-advanced-regression-techniques/submit>
- Upload “predict1.csv” file created above.
- Upload your score as an image below.

```
In [19]: from IPython.display import Image
Image(filename='images/asgn_01.png', width=500)
```

Out[19]:

Overview	Data	Kernels	Discussion	Leaderboard	More	My Submissions	Submit Predictions
Submission and Description		Private Score	Public Score	Use for Final Score			
predict3.csv a few seconds ago by Mudassir Khan replace negative by zero			0.49101	<input type="checkbox"/>			
predict3.csv 26 minutes ago by Mudassir Khan after normalization			6.33587	<input type="checkbox"/>			
predict3.csv an hour ago by Mudassir Khan using multivariate regression			Error ⓘ	<input type="checkbox"/>			
predict3.csv an hour ago by Mudassir Khan using more features			Error ⓘ	<input type="checkbox"/>			
predict2.csv an hour ago by Mudassir Khan using scikit learn			0.28918	<input type="checkbox"/>			
predict.csv 7 hours ago by Mudassir Khan			0.28719	<input type="checkbox"/>			

Task 3: Use scikit-learn for Linear Regression

In this task we are going to use [Linear Regression class from scikit-learn](#) library to train the same model. The aim is to move from understanding algorithm to using an existing well established library. There is a [Linear Regression example](#) available on scikit-learn website as well.

- Use the scikit-learn linear regression class to train the model on `df_train`
- Compare the parameters from scikit-learn `linear_model.LinearRegression.coef_` to the θ_s from earlier.
- Use the `linear_model.LinearRegression.predict` on test data and upload it to kaggle. See if your score improves. Provide screenshot.
- Note: no need to append 1's to `X_train`. Scikit linear regression has parameter called `fit_intercept` that is by default enabled.

```
In [20]: # import scikit-learn linear model
         from sklearn import linear_model

         # get X and y
         # write code here

         # Create linear regression object
```

```

# write code here check link above for example

# Train the model using the training sets. Use fit(X,y) command
# write code here

# The coefficients
print('Intercept: \n', regr.intercept_)
print('Coefficients: \n', regr.coef_)
# The mean squared error
print("Mean squared error: %.2f"
      % np.mean((regr.predict(X_train) - y_train) ** 2))
# Explained variance score: 1 is perfect prediction
print('Variance score: %.2f' % regr.score(X_train, y_train))

```

```

Intercept:
[ 18569.02585649]
Coefficients:
[[ 107.13035897]]
Mean squared error: 3139843209.67
Variance score: 0.50

```

```

In [21]: # read test X without 1's
# write code here

```

```

In [22]: # predict output for test data. Use predict(X) command.
predict2 = # write code here

```

```

In [23]: # remove negative sales by replacing them with zeros
predict2[predict2<0] = 0

```

```

In [24]: # save prediction as predict2.csv file
# write code here

```

Task 4: Multivariate Linear Regression

Lastly use columns ['OverallQual', 'GrLivArea', 'GarageCars', 'TotalBsmtSF', 'YearBuilt'] and scikit-learn or the code given above to predict output on test data. Upload it to kaggle like earlier and see how much it improves your score.

- Everything remains same except dimensions of X changes.
- There might be some data missing from the test or train data that you can check using pandas.DataFrame.describe() function. Below we provide some helping functions for removing that data.

```

In [25]: # define columns ['OverallQual', 'GrLivArea', 'GarageCars', 'TotalBsmtSF',
# write code here

```

```

# check features range and statistics. Training dataset looks fine as all
df_train[cols].describe()

```

```
Out [25]:
```

	OverallQual	GrLivArea	GarageCars	TotalBsmtSF	YearBuilt
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000
mean	6.099315	1515.463699	1.767123	1057.429452	1971.267808
std	1.382997	525.480383	0.747315	438.705324	30.202904
min	1.000000	334.000000	0.000000	0.000000	1872.000000
25%	5.000000	1129.500000	1.000000	795.750000	1954.000000
50%	6.000000	1464.000000	2.000000	991.500000	1973.000000
75%	7.000000	1776.750000	2.000000	1298.250000	2000.000000
max	10.000000	5642.000000	4.000000	6110.000000	2010.000000

```
In [26]: # Load X and y variables from pandas dataframe df_train
         # write code here
```

```
         # Get m = number of samples and n = number of features
         # write code here
```

```
In [27]: #Feature normalizing the columns (subtract mean, divide by standard deviation)
         #Store the mean and std for later use
         #Note don't modify the original X matrix, use a copy
         stored_feature_means, stored_feature_stds = [], []
         Xnorm = np.array(X_train).copy()
         for icol in range(Xnorm.shape[1]):
             stored_feature_means.append(np.mean(Xnorm[:,icol]))
             stored_feature_stds.append(np.std(Xnorm[:,icol]))
             #Skip the first column if 1's
             # if not icol: continue
             #Faster to not recompute the mean and std again, just used stored values
             Xnorm[:,icol] = (Xnorm[:,icol] - stored_feature_means[-1])/stored_feature_stds[-1]

         # check data after normalization
         pd.DataFrame(data=Xnorm, columns=cols).describe()
```

```
Out [27]:
```

	OverallQual	GrLivArea	GarageCars	TotalBsmtSF	YearBuilt
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000
mean	0.083562	0.036301	-0.232877	0.069863	0.013014
std	0.655971	0.732612	0.747315	0.717391	0.734330
min	-3.000000	-2.000000	-2.000000	-2.000000	-3.000000
25%	0.000000	0.000000	-1.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000	0.000000
max	2.000000	7.000000	2.000000	11.000000	1.000000

```
In [28]: # Run Linear Regression from scikit-learn or code given above.

         # write code here. Repeat from above.
```

```
Intercept:
[-672030.53248541]
Coefficients:
```

```
[[ 20391.14093374      50.83150559  14510.0032998      29.97787732
    301.43341059]]
Mean squared error: 1463421280.61
Variance score: 0.77
```

```
In [29]: # To predict output using ['OverallQual', 'GrLivArea', 'GarageCars', 'TotalBsmtSF']
# Check features range and statistics to see if there is any missing data.
# As you can see from count "GarageCars" and "TotalBsmtSF" has 1 missing value
df_test[cols].describe()
```

```
Out[29]:
```

	OverallQual	GrLivArea	GarageCars	TotalBsmtSF	YearBuilt
count	1459.000000	1459.000000	1458.000000	1458.000000	1459.000000
mean	6.078821	1486.045922	1.766118	1046.117970	1971.357779
std	1.436812	485.566099	0.775945	442.898624	30.390071
min	1.000000	407.000000	0.000000	0.000000	1879.000000
25%	5.000000	1117.500000	1.000000	784.000000	1953.000000
50%	6.000000	1432.000000	2.000000	988.000000	1973.000000
75%	7.000000	1721.000000	2.000000	1305.000000	2001.000000
max	10.000000	5095.000000	5.000000	5095.000000	2010.000000

```
In [30]: # Replace missing value with the mean of the feature
df_test['GarageCars'] = df_test['GarageCars'].fillna((df_test['GarageCars'].mean()))
df_test['TotalBsmtSF'] = df_test['TotalBsmtSF'].fillna((df_test['TotalBsmtSF'].mean()))
```

```
In [31]: df_test[cols].describe()
```

```
Out[31]:
```

	OverallQual	GrLivArea	GarageCars	TotalBsmtSF	YearBuilt
count	1459.000000	1459.000000	1459.000000	1459.000000	1459.000000
mean	6.078821	1486.045922	1.766118	1046.117970	1971.357779
std	1.436812	485.566099	0.775679	442.746712	30.390071
min	1.000000	407.000000	0.000000	0.000000	1879.000000
25%	5.000000	1117.500000	1.000000	784.000000	1953.000000
50%	6.000000	1432.000000	2.000000	988.000000	1973.000000
75%	7.000000	1721.000000	2.000000	1304.000000	2001.000000
max	10.000000	5095.000000	5.000000	5095.000000	2010.000000

```
In [32]: # read test X without 1's
# write code here
```

```
# predict using trained model
predict3 = # write code here
```

```
# replace any negative predicted saleprice by zero
predict3[predict3<0] = 0
```

```
In [33]: # predict target/output variable for test data using the trained model and
# write code to save output as predict3.csv here
```

5 Resources

Course website: <https://w4zir.github.io/ml17s/>
[Course resources](#)

6 Credits

Raschka, Sebastian. Python machine learning. Birmingham, UK: Packt Publishing, 2015. Print.
Andrew Ng, Machine Learning, Coursera
[Scikit Learn Linear Regression](#)